

Chapter 1

An Introduction to PyTorch

1.1 A Fun Example

Efficient machine learning processes data in batches, and our model will expect a batch of data.

We use PyTorch's `unsqueeze()` function to add a dimension to our tensor and create a batch of size 1.

The use of `model.to(device)` and `batch.to(device)` sends our model and input data to the GPU if available, and executing `model(batch.to(device))` runs our classifier.

Chapter 2

Tensors

2.1 Creating Tensors

Use `torch.arange()` when the step size is known. Use `torch.linspace()` when the number of elements is known. You can use `torch.tensor()` to create tensors from array-like structures such as lists, NumPy arrays, tuples, and sets. To convert existing tensors to NumPy arrays and lists, use the `torch.numpy()` and `torch.tolist()` functions, respectively.

Table 2.1: Tensor creation functions

Function	Description
<code>torch.tensor(data, dtype=None, device=None, requires_grad=False, pin_memory=False)</code>	Creates a tensor from an existing data structure
<code>torch.empty(*size, out=None, dtype=None, layout=torch.strided, device=None, requires_grad=False)</code>	Creates a tensor from uninitialized elements based on the random state of values in memory
<code>torch.zeros(*size, out=None, dtype=None, layout=torch.strided, device=None, requires_grad=False)</code>	Creates a tensor with all elements initialized to 0.0
<code>torch.ones(*size, out=None, dtype=None, layout=torch.strided, device=None, requires_grad=False)</code>	Creates a tensor with all elements initialized to 1.0
<code>torch.arange(start=0, end, step=1, out=None, dtype=None, layout=torch.strided, device=None, requires_grad=False)</code>	Creates a 1D tensor of values over a range with a common step value
<code>torch.linspace(start, end, steps=100, out=None, dtype=None, layout=torch.strided, device=None, requires_grad=False)</code>	Creates a 1D tensor of linearly spaced points between the start and end
<code>torch.logspace(start, end, steps=100, base=10.0, out=None, dtype=None, layout=torch.strided, device=None, requires_grad=False)</code>	Creates a 1D tensor of logarithmically spaced points between the start and end
<code>torch.eye(n, m=None, out=None, dtype=None, layout=torch.strided, device=None, requires_grad=False)</code>	Creates a 2D tensor with ones on the diagonal and zeros everywhere else
<code>torch.full(size, fill_value, out=None, dtype=None, layout=torch.strided, device=None, requires_grad=False)</code>	Creates a tensor filled with fill_value
<code>torch.load(f)</code>	Loads a tensor from a serialized pickle file
<code>torch.save(f)</code>	Saves a tensor to a serialized pickle file